
Learning with kernels and SVMs

Théomé Borck, February 8th, 2021

Seminar “Mathematics for Data Science”



References

- Schölkopf, B and Smola, A. (2002) [Learning with kernels, Support Vector Machines, Regularization, Optimization and Beyond](#). MIT Press, Cambridge, MA.
- Poggio T. and Smale, S. (2003) [The Mathematics of Learning: Dealing with Data](#). *Notices of the American Mathematical Society (AMS)*, Vol. 50, No. 5, 537-544.

Additional references

- Zisserman, A. (2015) [Lecture 3: SVM dual, kernels and regression](#). *C19 Machine Learning*, Oxford University, Oxford.
- Günnemann, S. (2019) [Lecture 8: SVM and kernels](#). *IN2064 Machine Learning WS2019*, TUM, Munich.

Table of Contents

1. Support Vector Machine

SVM and soft-SVM, primal and dual problems, support vectors

2. Kernels

Feature map projection, the “kernel trick”, Mercer’s theorem

3. SVMs in practice

Multi-classification, applications, support vector regression


4. The mathematics of learning

Dealing with data, a theoretical approach to bias/variance tradeoff and its link with SVM


Introduction

Binary classification


= given training data, learn a **classifier** to discriminate data instances



Linearly separable



Non-linearly separable



$$(x_i, y_i), i \in [N], y_i \in \{-1, 1\} \quad \rightarrow \quad \begin{aligned} f(x_i) &= w^t x_i + b \\ y_i f(x_i) &> 0 \end{aligned}$$

→ **Perceptron algorithm**

Converges slowly, only works with linearly-separable data and no margin interpretation

1. Support Vector Machine

The intuition


Objective = Find a hyperplane that separates both classes with the **maximum margin**

→ more stable model given perturbations of the training data

→ better **generalization**

Actual rigorous motivation

→ Statistical Learning Theory (Vapnik, 1995)



1. Support Vector Machine

An optimization problem

The SVM objective comes down to:

$$\text{Minimize } f_0(\mathbf{w}, b) = \frac{1}{2} \mathbf{w}^T \mathbf{w} \quad \leftarrow \text{L2-norm}$$


$$\text{subject to } f_i(\mathbf{w}, b) = y_i(\mathbf{w}^T \mathbf{x}_i + b) - 1 \geq 0 \quad \text{for } i = 1, \dots, N.$$

This is a **convex** and **quadratic** optimization problem, subject to **linear constraints**.


→ a global optimal solution can be obtained, i.e. a unique minimizer.

1. Support Vector Machine

Soft-SVM: introducing slack variables



Points can be linearly separated but the margin is very **narrow**.



The larger margin solution is better, but **one constraint is violated**.



We introduce a “**trade-off**” between the margin and the number of mistakes on the training data

1. Support Vector Machine

Soft-SVM: introducing slack variables

Idea = Relax the constraints but punish as much as necessary the relaxation of a constraint


→ slack variable $\xi_i \geq 0$

$$y_i(\mathbf{w}^T \mathbf{x}_i + b) \geq 1 - \xi_i \quad \text{for all } i.$$

For $0 < \xi_i \leq 1$, the instance is between margin and correct side of the hyperplane.

→ margin violation

For $\xi_i > 1$, point is misclassified.



1. Support Vector Machine

Soft-SVM: introducing slack variables

The SVM objective becomes:

$$\begin{aligned} \text{Minimize} \quad & f_0(\mathbf{w}, b, \boldsymbol{\xi}) = \frac{1}{2} \mathbf{w}^T \mathbf{w} + C \sum_{i=1}^N \xi_i \\ \text{subject to} \quad & y_i(\mathbf{w}^T \mathbf{x}_i + b) - 1 + \xi_i \geq 0 & i = 1, \dots, N, \\ & \xi_i \geq 0 & i = 1, \dots, N. \end{aligned}$$

$C > 0$ determines how heavy is the violation punished \rightarrow [regularization](#) parameter.

The best C is often found through [hyperparameter search](#).

1. Support Vector Machine

SGD to solve the SVM problem - Pegasos algorithm

The slack variables can be concisely rewritten as:


$$\xi_i = \begin{cases} 1 - y_i(\mathbf{w}^T \mathbf{x}_i + b), & \text{if } y_i(\mathbf{w}^T \mathbf{x}_i + b) < 1 \\ 0 & \text{else} \end{cases}$$

Hence, the problem becomes an unconstrained optimization problem:

$$\min_{\mathbf{w}, b} \frac{1}{2} \mathbf{w}^T \mathbf{w} + C \sum_{i=1}^N \max\{0, 1 - y_i(\mathbf{w}^T \mathbf{x}_i + b)\}$$

Regularization

Hinge loss



→ can be optimized using standard gradient-based methods such as [Stochastic Gradient Descent](#).

1. Support Vector Machine

From primal to dual problem

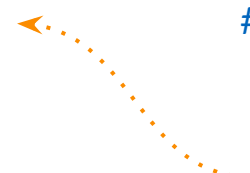
Using the **Lagrangian multipliers** method (or through the **Representer Theorem**), we can convert our problem:

Primal problem

$$\begin{aligned}
 \text{Minimize} \quad & f_0(\mathbf{w}, b, \boldsymbol{\xi}) = \frac{1}{2} \mathbf{w}^T \mathbf{w} + C \sum_{i=1}^N \xi_i && \# \text{ parameters} = \text{vector space dim.} = d \\
 \text{subject to} \quad & y_i(\mathbf{w}^T \mathbf{x}_i + b) - 1 + \xi_i \geq 0 && i = 1, \dots, N, \\
 & \xi_i \geq 0 && i = 1, \dots, N.
 \end{aligned}$$


Dual problem

$$\begin{aligned}
 \text{Maximize} \quad & g(\boldsymbol{\alpha}) = \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N y_i y_j \alpha_i \alpha_j \mathbf{x}_i^T \mathbf{x}_j && \# \text{ parameters} = N \\
 \text{subject to} \quad & \sum_{i=1}^N \alpha_i y_i = 0 && 0 \leq \alpha_i \leq C
 \end{aligned}$$


 Inner product of data instances

1. Support Vector Machine

Support vectors




1. Support Vector Machine


Where do we stand ?

- We have a **linear binary classifier** which **maximizes the margin** between two classes.
- Our objective function is convex so we can **optimize using gradient descent**.
- We derived the SVM through a dual formulation, solvable using standard QP libraries.

But how do we handle the more complex non-linearly separable cases ?



Handled by slack variables



Underlying non-linear distribution

2. Kernels

Feature map projection

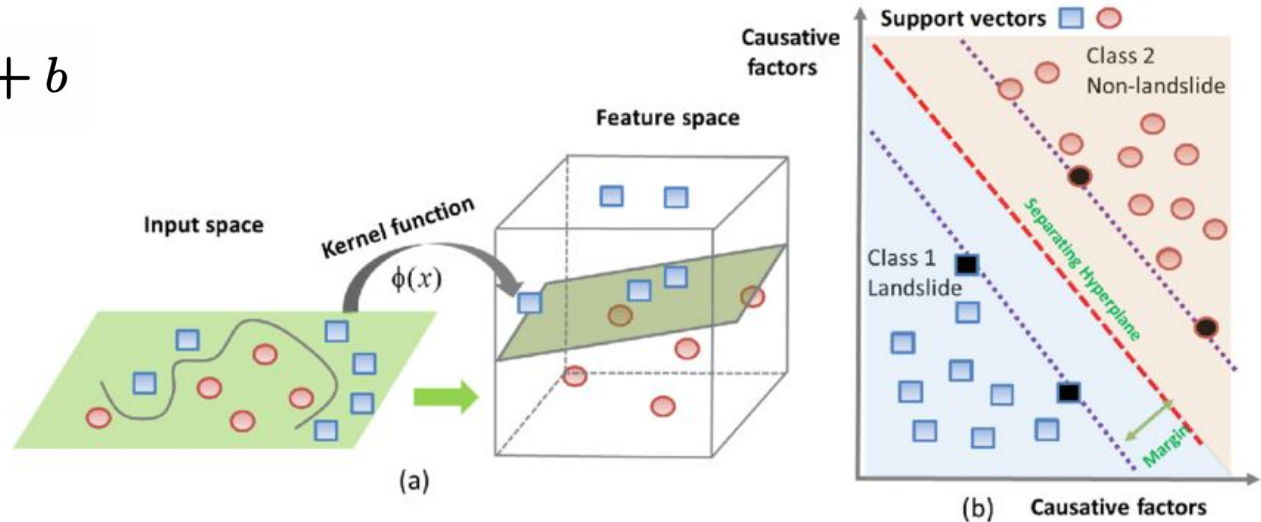
To find a linearly-separable space, we project to a **higher-dimensional space**.

$$\Phi : \mathbf{x} \rightarrow \Phi(\mathbf{x}) \quad \mathbb{R}^d \rightarrow \mathbb{R}^D$$

The linear classifier to learn is now in this projected space.

$$f(\mathbf{x}) = \mathbf{w}^\top \Phi(\mathbf{x}) + b$$

Feature map



2. Kernels

The “kernel trick”

With the feature projection, our dual formulation

$$g(\boldsymbol{\alpha}) = \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N y_i y_j \alpha_i \alpha_j \mathbf{x}_i^T \mathbf{x}_j \quad \text{becomes} \quad g(\boldsymbol{\alpha}) = \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N y_i y_j \alpha_i \alpha_j \boldsymbol{\phi}(\mathbf{x}_i)^T \boldsymbol{\phi}(\mathbf{x}_j)$$

Defining the **kernel function**: $k : R^d \times R^d \rightarrow R$

$$k(\mathbf{x}_i, \mathbf{x}_j) := \boldsymbol{\phi}(\mathbf{x}_i)^T \boldsymbol{\phi}(\mathbf{x}_j)$$

We can rewrite the dual as

$$g(\boldsymbol{\alpha}) = \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N y_i y_j \alpha_i \alpha_j k(\mathbf{x}_i, \mathbf{x}_j)$$



The classifier can be learnt and applied **without explicitly computing** the feature map projection. Complexity of learning only depends on N and not on D.

2. Kernels

Examples


Linear kernel: $k(\mathbf{a}, \mathbf{b}) = \mathbf{a}^T \mathbf{b}$

Polynomial kernel $k(\mathbf{a}, \mathbf{b}) = (\mathbf{a}^T \mathbf{b})^p$ or $(\mathbf{a}^T \mathbf{b} + 1)^p$ >

→ contains all polynomials terms up to degree p

Gaussian kernel: $k(\mathbf{a}, \mathbf{b}) = \exp\left(-\frac{\|\mathbf{a} - \mathbf{b}\|^2}{2\sigma^2}\right)$ >

→ projection to an infinite dimensional feature space



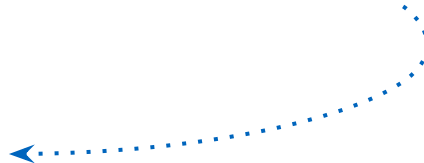
$$\begin{aligned} \text{L} \rightarrow \exp\left(-\frac{1}{2} \|x - y\|_2^2\right) &= \underbrace{\exp\left(-\frac{1}{2} \|x\|_2^2\right)}_{c(x)} \underbrace{\exp\left(-\frac{1}{2} \|y\|_2^2\right)}_{c(y)} \sum_{k=0}^{\infty} \frac{(x, y)^k}{k!} = \sum_{k=0}^{\infty} \left\langle \sqrt[k]{\frac{c(x)}{k!}} x, \sqrt[k]{\frac{c(y)}{k!}} y \right\rangle^k \end{aligned}$$

2. Kernels

Extending kernels to non-numerical data

Mercer's theorem:


A kernel is valid if it gives rise to a **symmetric, positive semidefinite kernel matrix \mathbf{K}** for any input data \mathbf{X} .

$$\mathbf{K} = \begin{pmatrix} k(\mathbf{x}_1, \mathbf{x}_1) & k(\mathbf{x}_1, \mathbf{x}_2) & \cdots & k(\mathbf{x}_1, \mathbf{x}_N) \\ k(\mathbf{x}_2, \mathbf{x}_1) & k(\mathbf{x}_2, \mathbf{x}_2) & \cdots & k(\mathbf{x}_2, \mathbf{x}_N) \\ \vdots & \vdots & \ddots & \vdots \\ k(\mathbf{x}_N, \mathbf{x}_1) & k(\mathbf{x}_N, \mathbf{x}_2) & \cdots & k(\mathbf{x}_N, \mathbf{x}_N) \end{pmatrix}$$


Hence, we can use kernels to **encode similarity with non-numerical data** such as strings or graphs.

$$\rightarrow k(\text{"Coca"}, \text{"Pepsi"}) = -5 \quad k(\text{"Coca"}, \text{"Wine"}) = 8$$

$$f(\mathbf{x}) = \text{sgn} \left(\sum \lambda_i k(\mathbf{x}, \mathbf{x}_i) + b \right)$$



Classification

Weights

Comparison $\rightarrow K(X, X_i)$

Support vectors

Input data

3. SVM in practice

SVM for multi-classification

Two versions:

One-vs-rest = Train C SVM models for C classes, where each SVM is being trained for classification of one class against all the remaining ones.

→ The predicted class is then the one, where the distance from the hyperplane is maximal.

One-vs-one

= Train a classifier for all possible pairings and evaluate all.

→ The predicted class is the one with the majority vote
(the votes are weighted according to the distance from the margin).

3. SVM in practice

Applications

- Face detection
- Image classification
- Handwriting recognition
- Bioinformatics - protein/gene classification
- News classification
- ... and many other classification problems.

MNIST Benchmark

60000 training examples

10000 test examples

28x28



Classifier	test error	reference
linear classifier	8.4%	[7]
3-nearest-neighbour	2.4%	[7]
SVM	1.4%	[12]
Tangent distance	1.1%	[57]
LeNet4	1.1%	[38]
Boosted LeNet4	0.7%	[38]
Translation invariant SVM	0.56%	[18]

3. SVM in practice

SVM for regression


We want to find a **prediction line** (or an hyperspace in a high-dimensional feature space) which deviates less than ε with respect to the ground truth y .

$$\min_{w, \xi_i} \underbrace{\frac{1}{2} \|w\|_2^2}_{\text{Regularization}} + C \sum_{i=1}^N |\xi_i|$$

Regularization

subject to

$$|y_i - w^t x_i| \leq \varepsilon + |\xi_i|$$



4. The mathematics of learning

A key algorithm to learning theory

“The Mathematics of Learning: Dealing with Data.” from Poggio and Smale, written in 2003.

→ Outlines the mathematical foundations of **learning theory** and describe a key algorithm of it.

Supervised learning = systems trained, instead of programmed, with a set of examples.

Training = synthesize a function that best represents the relation between inputs and outputs.

Learning = a principle method for distilling predictive and hence scientific theories from the data.

4. The mathematics of learning

A key algorithm to learning theory

1. Start with data $(x_i, y_i)_{i=1}^m$
2. Choose a symmetric, positive definite function $K_x(x') = K(x, x')$, continuous on $X \times X$. A kernel $K(t, s)$ is *positive definite* if $\sum_{i,j=1}^n c_i c_j K(t_i, t_j) \geq 0$ for any $n \in \mathbb{N}$ and choice of $t_1, \dots, t_n \in X$ and $c_1, \dots, c_n \in \mathbb{R}$. An example of such a Mercer kernel is the Gaussian

$$K(x, x') = e^{-\frac{\|x-x'\|^2}{2\sigma^2}}. \quad (1)$$

restricted to $X \times X$.

3. Set $f : X \rightarrow Y$ to

$$f(x) = \sum_{i=1}^m c_i K_{x_i}(x). \quad (2)$$

where $\mathbf{c} = (c_1, \dots, c_m)$ and

$$(m\gamma \mathbf{I} + \mathbf{K})\mathbf{c} = \mathbf{y} \quad (3)$$

where \mathbf{I} is the identity matrix, \mathbf{K} is the square positive definite matrix with elements $K_{i,j} = K(x_i, x_j)$ and \mathbf{y} is the vector with coordinates y_i . The parameter γ is a positive, real number.

→ Equation 2 approximates the unknown function by a weighted superposition of Gaussian “blobs”, each centered at the location of one of the m examples.

→ The algorithm performs well in a number of applications involving regression as well as binary classification.

4. The mathematics of learning

The derivation of this key algorithm

The algorithm can be derived from [Tikhonov regularization](#).

We use [Empirical Risk Minimization](#) to find a function f which minimizes:

$$\frac{1}{m} \sum_{i=1}^m (y_i - f(x_i))^2 + \gamma \|f\|_K^2,$$

Norm in \mathcal{H}_K - the [Reproducing Kernel Hilbert Space \(RKHS\)](#), defined by the kernel K

➤ [Proximal vector machines](#)

➤ For SVM, we replace the loss function by a more adapted non-quadratic loss $V(f(x, y)) = (1 - yf(x))_+$
Coefficients are then retrieved by solving the [quadratic programming problem](#).

➤ [Bayesian interpretation](#) of the regularization term

4. The mathematics of learning

A theoretical approach to the bias/variance trade-off

We can link the algorithm to some basics of the learning theory, which yields:

$$\int_X (f_z - f_\rho)^2 = S(z, \mathcal{H}) + \int_X (f_{\mathcal{H}} - f_\rho)^2 = S(z, \mathcal{H}) + A(\mathcal{H})$$

S must be estimated in probability over z and the estimate is called the **sample error**. It is studied thanks to the **theory of probability**.

A is dealt with via **approximation theory** and is called the **approximation error**.

- Related to the **bias/variance trade-off** in statistics.
- We can apply theoretical bounds to these errors.

4. The mathematics of learning

A theoretical framework to understand SVM and its derivation

Basis concept of SVM in linear separable case = *separation by an hyperplane maximizing the margin.*

In the non-separable case, this interpretation loses most of its meaning.

→ A more general and simpler framework for deriving SVM algorithms for classification and regression is to regard them as *special cases of regularization and follow the proposed key algorithm.*



Maximizing the margin is exactly equivalent to minimizing $\|w\|^2$ which corresponds to minimizing the RKHS norm.

Conclusion

Support Vector Machine

- A classification or regression algorithm which learns from labeled data → supervised.
- Can be interpreted through the concepts of hyperplanes and margin, but also more theoretically from the learning theory.

Kernels

- Provide a “trick” to help SVMs handle the non-linearly separable data.
- Also take root in a key algorithm derived from the learning theory.
- Choosing a kernel can be challenging → *no free lunch ;)*

Thank you!
Any questions ?